

How much maths do I need to learn to be a data scientist?

Outcomes for this evening

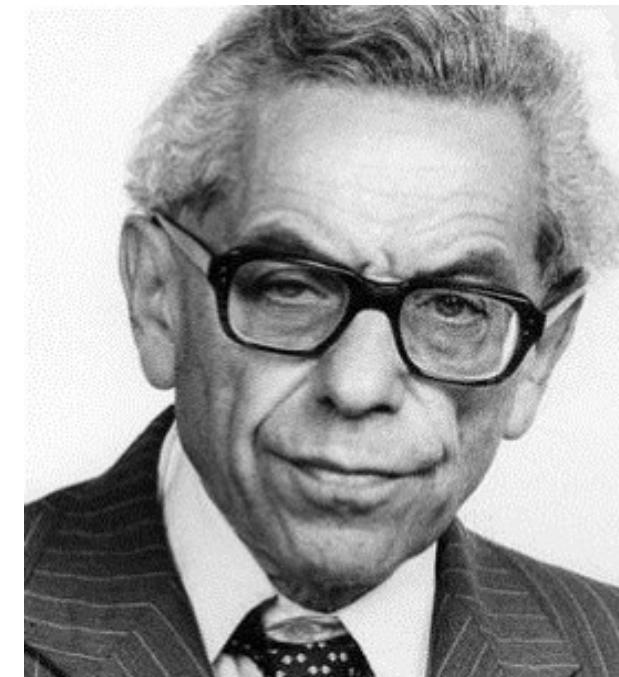
By the end of today's class, you should be able to answer the following:

- Randomly select elements from a list
- Explain the role of sampling data
- Transpose a matrix and a dataframe
- Explain uses of Gaussian and uniform distributions
- Identify misleading representations of data
- Describe three biases that can occur when gathering data

- Math and Data Science
- Probability
- Correlation
- Sets
- Linear Algebra and Numpy
- Calculus and Differential Equations
- Visualization
- Homework

Consider this a survey of an old, diverse field

- [Linear algebra](#) (vectors, matrices, cross product)
- Sets (union)
- [Statistics](#) (ie mean, median)
 - [How to Lie with Statistics](#)
 - Error bars
- [Probability](#), combinatorics
- Calculus
- ODEs and PDEs



*"I hope we'll be
able to solve
these problems
before we leave."*

Paul Erdős

Caveat: I'm not a mathematician

I won't be able to teach you all of Math



Jargon | concept | example

Where Math shows up in Data Science

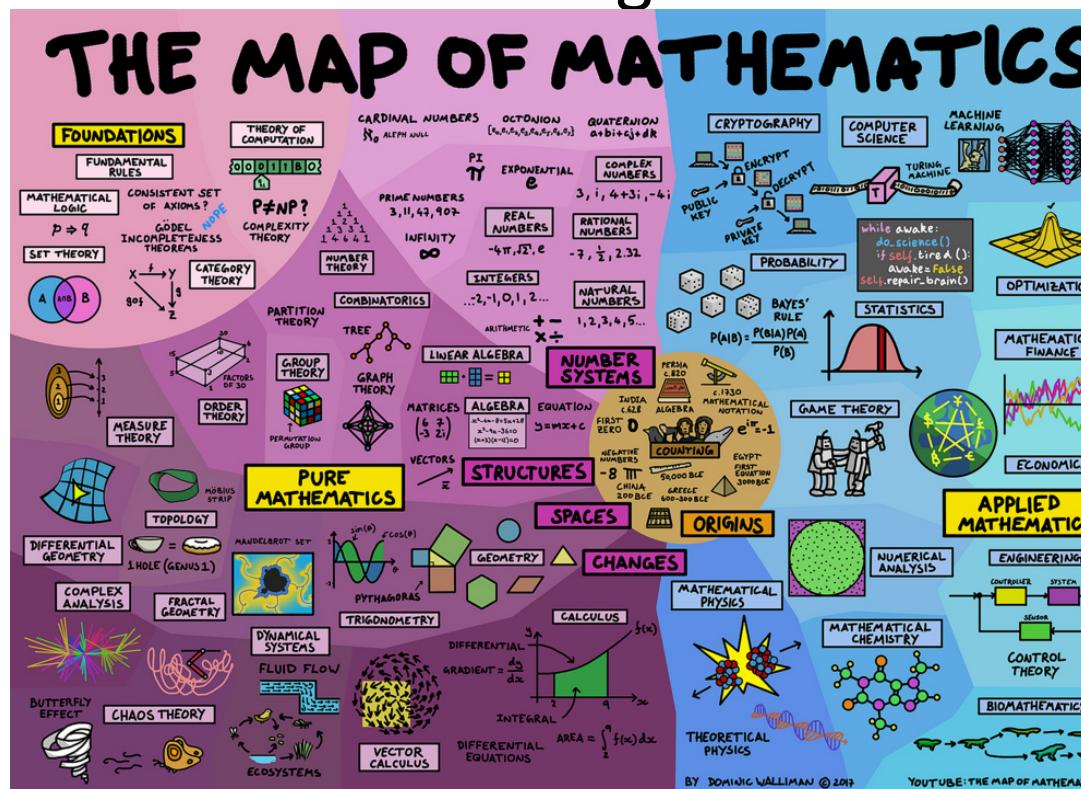
- Cleaning data – *filling in gaps for missing data with interpolation*
- Modeling expectations – *sense-making, distribution of each variable*
- Generating hypotheses that are numerically testable
- Evaluating test results to validate hypotheses
- Analysis of results – *visualization, sanity checks*
- Explanation of story to audience – what do you expect customers to take away? What is their language?

Resources for learning Math

- Focus on learning the jargon; this is necessary for searching
- For a given topic, evaluate the many options before investing time
 - Teaching style
 - Level of complexity
 - Assumptions about you, the student reader

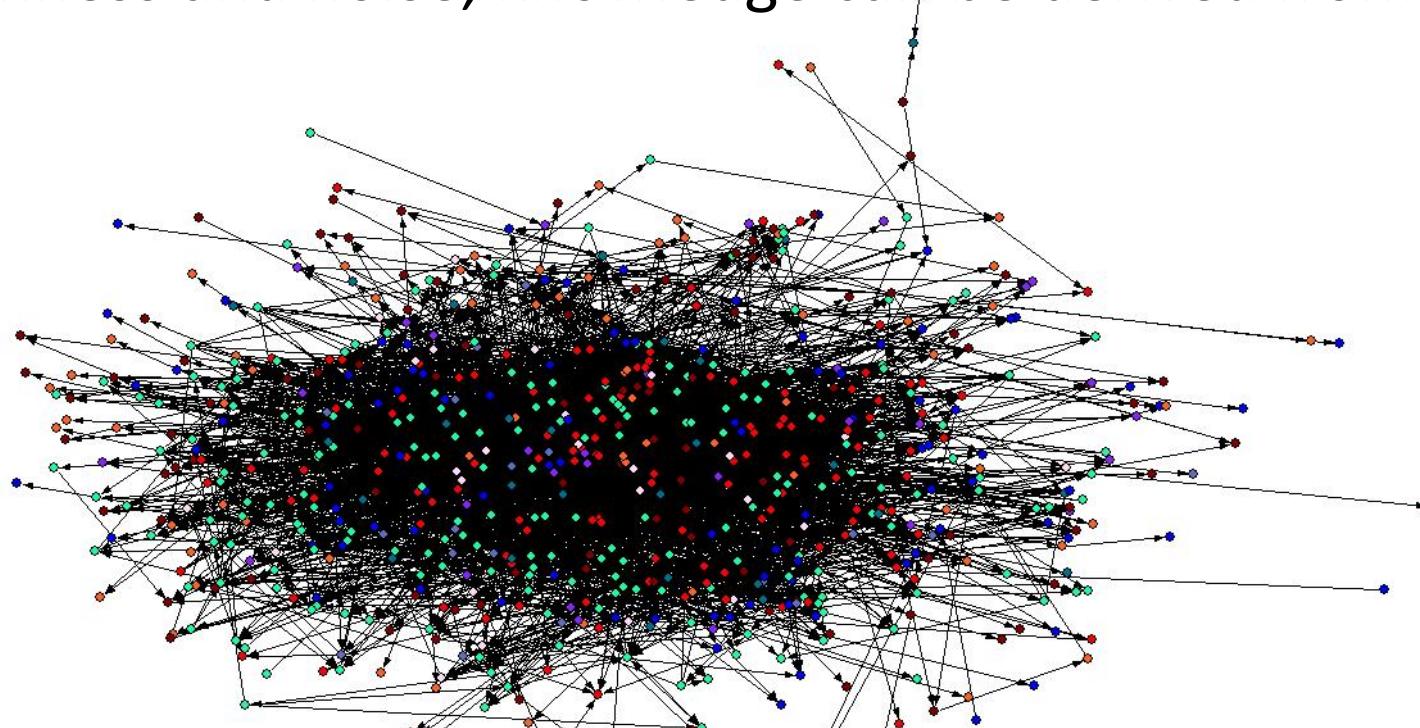
Free resources

- Online (blogs, [Coursera](#), YouTube)
- Books



Relevance of Statistics in Data Science

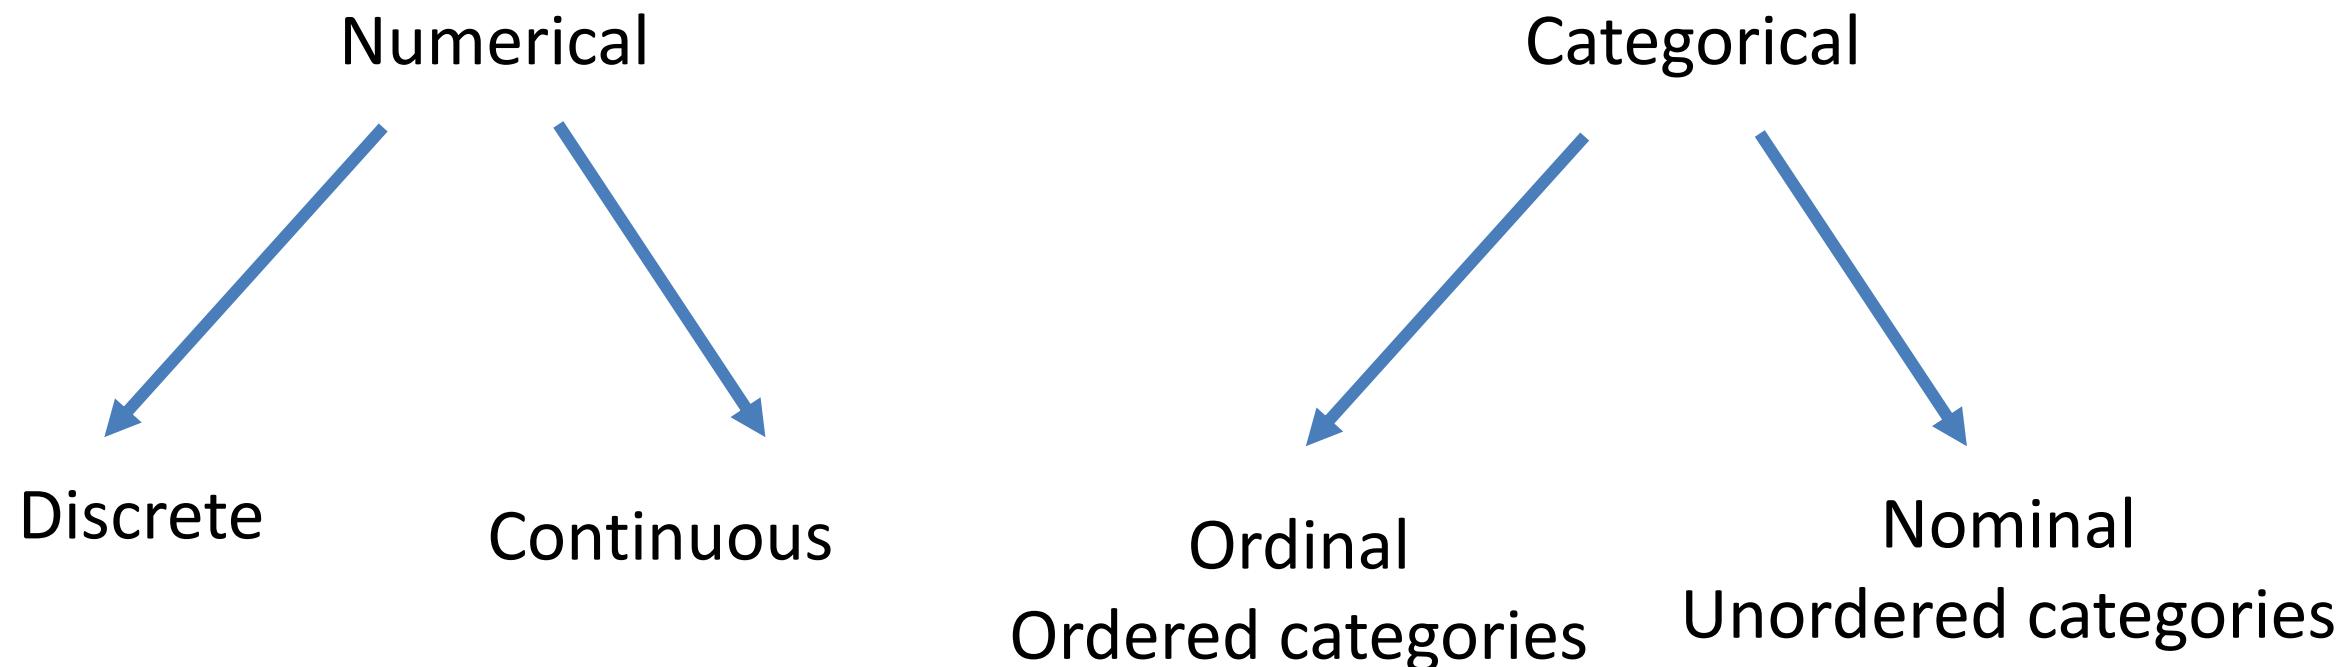
Given randomness and noise, knowledge can be derived from complex data



How:

Quantify relationships between variables in a model using standard language and techniques

Quantifying Relations



Discrete versus Continuous variables

- Discrete variables: outcomes for coin flips, deck of cards, roll of dice
- Continuous variables: time, distance

Continuous: For any two values of a variable, it is possible to get a measurement that is between the two values.

Trick: Rounding continuous to discrete

Rounding is often applied to continuous to make the variable discrete



- ~~Math and Data Science~~

- Probability
- Correlation
- Sets
- Linear Algebra and Numpy
- Calculus and Differential Equations
- Visualization
- Homework

Core to Statistics: Probability

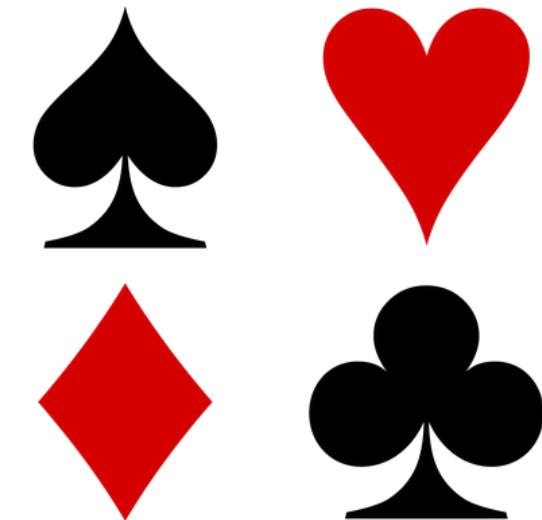
Statistics quantifies relationships
between variables in a model
using standard language and techniques

Probability is a way of figuring out an applicable model

Uniform distribution

Each outcome is equally likely:

Chance of any of the possible options is the same as any other outcome



Given a standard deck of 52 cards, what is the chance of getting a card that is a diamond?

Chance of getting a heart: 13/52

Chance of getting a diamond: 13/52

Chance of getting a club: 13/52

Chance of getting a spade: 13/52



Randomness and random selection in Python

getting_started_with_random.ipynb

```
random.choice( ['red', 'black', 'green', 5] )
```

Activity: Coin toss

Using your penny, write down (in order) tails/heads for 10 flips



How many possible permutations?

- 1 flip has 2 outcomes: head (H) or tails (T)
- 2 flips --> 4 outcomes: HH or HT or TH or TT
- 3 flips --> 8 outcomes: HHH,HHT,HTH,THH,HTT,THT,TTH,TTT
- 4 flips --> 16 outcomes: HHHH,HHHT,HHTH,HTHH,THHH,HHTT,...
- ...
- For N flips there are 2^N outcomes
- N=10 flips: $2^{10}= 1,024$ outcomes
- N=20 flips: $2^{20}=1,048,576$ outcomes

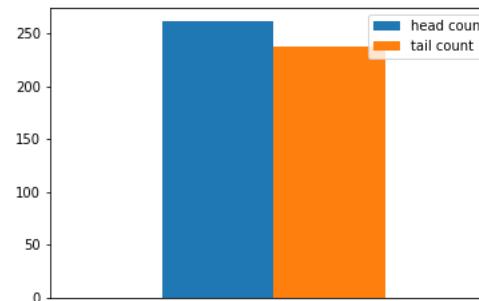
modeling_random_coin_flips.ipynb

Want to watch how this executes? Check out
<http://pythontutor.com/visualize.html>

Visualizing probabilistic outcomes

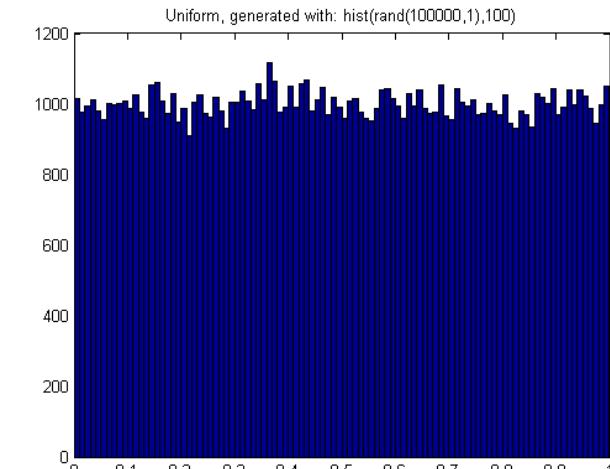
- A fair deck of cards has a uniform distribution of outcomes for a given selection
- A fair coin has a uniform distribution of outcomes

`modeling_random_coin_flips_visualization.ipynb`



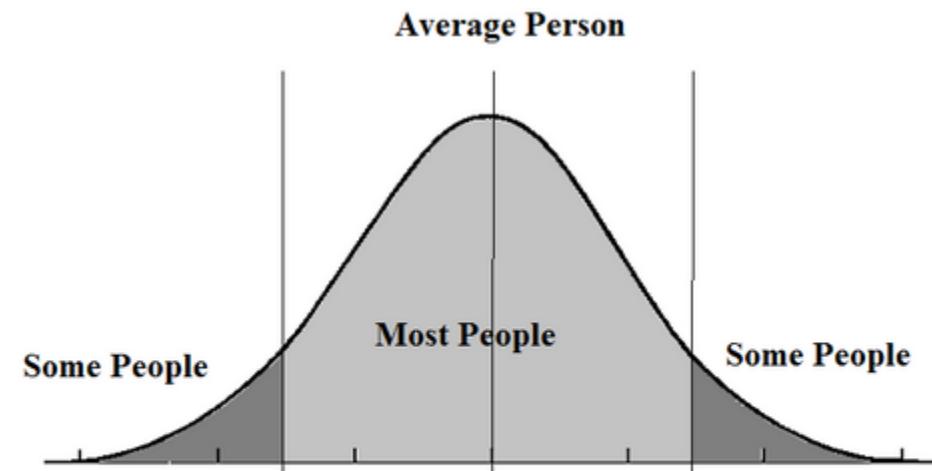
Two outcomes

Uniform distribution



many outcomes

Another distribution: the Bell curve



Gaussian and Binomial Distribution

[binomial distribution](#) is discrete; [normal](#) (Gaussian) is continuous

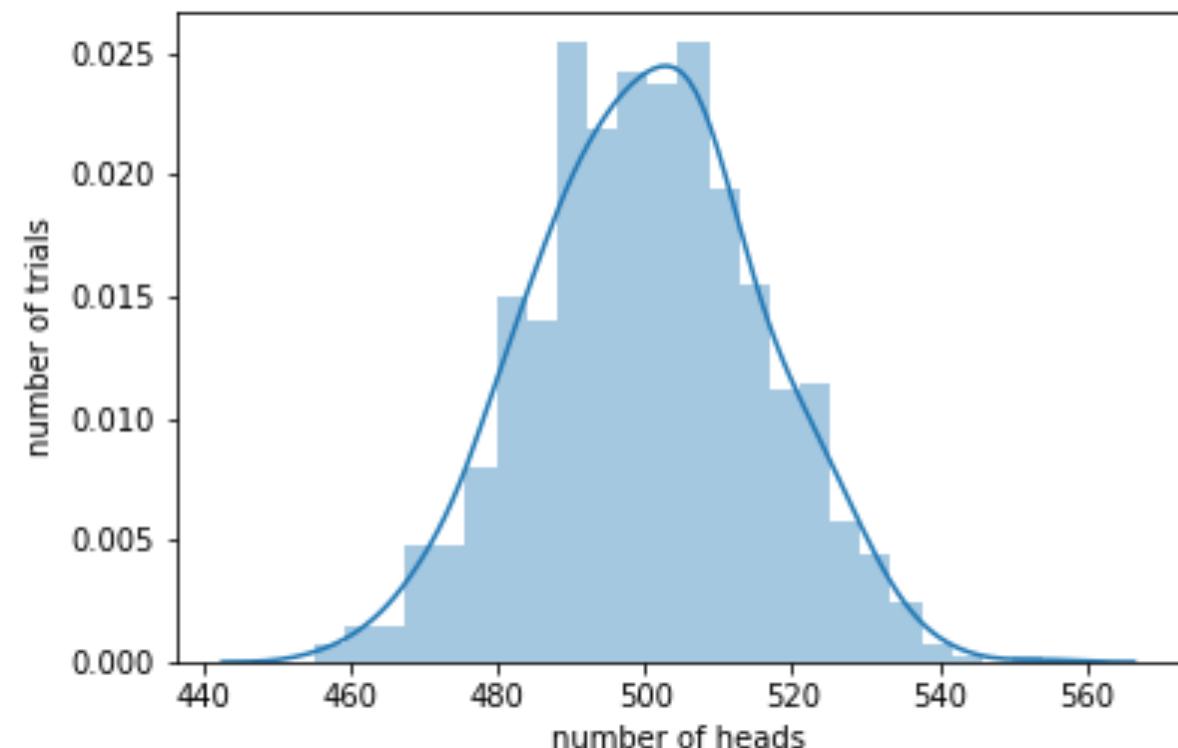
Characteristics of binomial variable:

- A fixed number of repeated, identical, independent trials. n is usually the parameter chosen to label the number of trials.
- Every trial results in either a success, with probability p , or a failure, with probability $1-p$. These must be the only two possible outcomes for a trial.
- The random variable of interest is the total number of trials that ended in a success.

Coin flips produce a bell curve!

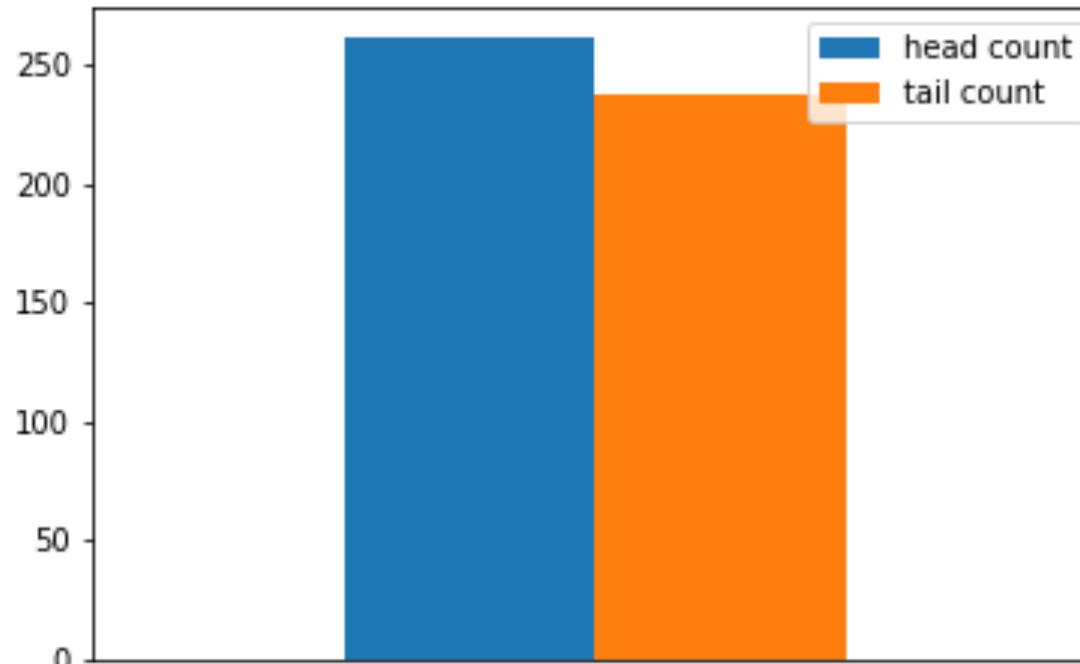
An exploration of the coin flips:

`binomial_distribution_for_coin_flips.ipynb`



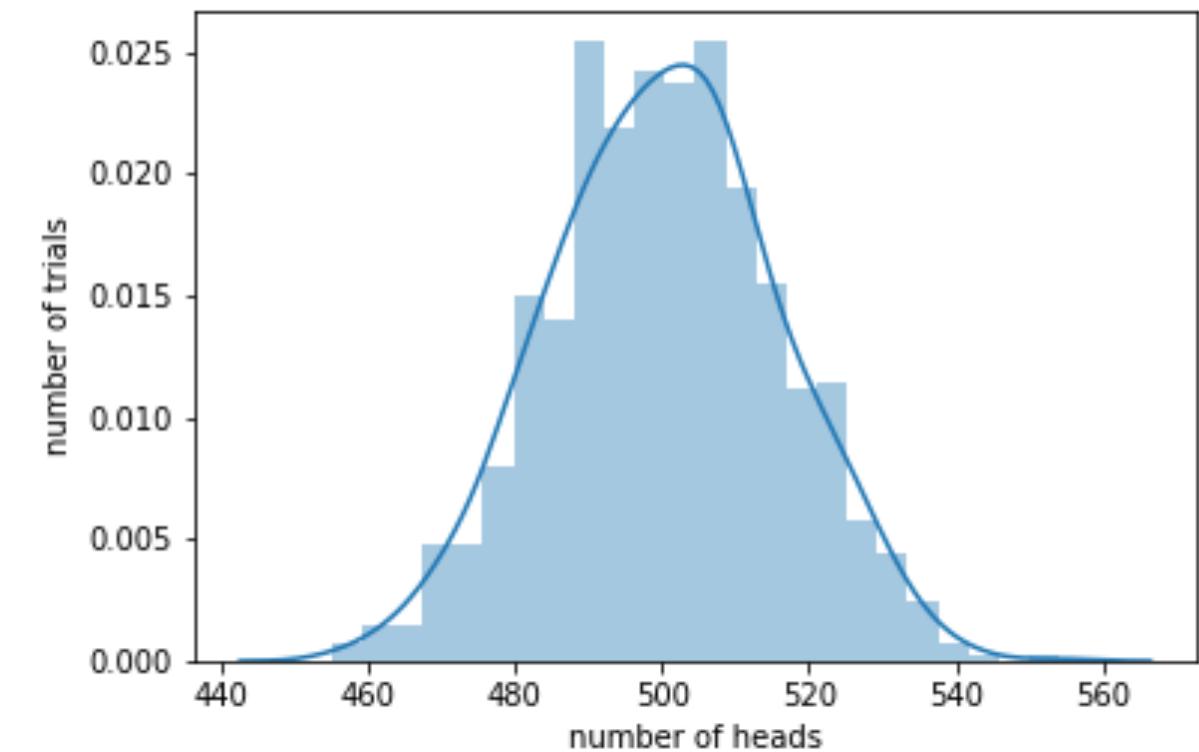
What is the relation between the curves?

Uniform distribution for single trial



Single experiment

Binomial distribution for many trials



Many experiments

Results vary: Error bars tell by how much

- Confidence Interval = certainty of what the mean value is
Confidence interval improves when more data is added

Distinct from

- Variance measures the width of a distribution
- Standard deviation is $\text{sqrt}(\text{variance})$ and has same units as variable

Variance and Standard Deviation do not change as the population size increases

Visually include info about distribution of variable in Violin plots

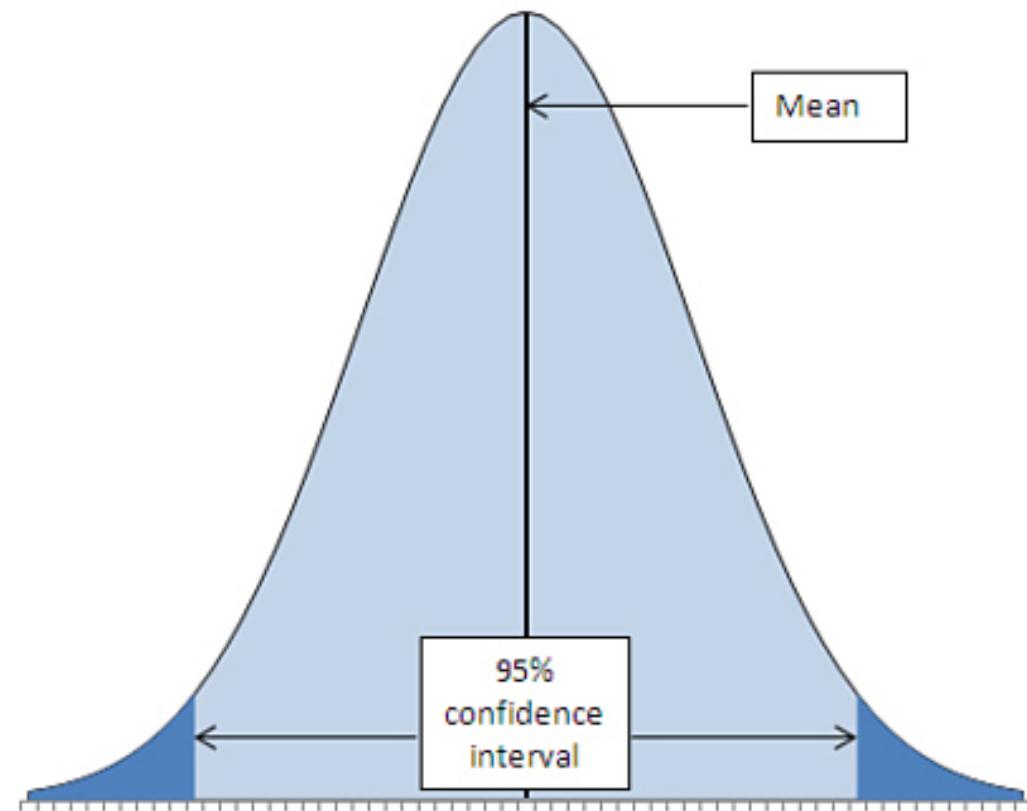
Results vary: Error bars tell by how much

- Confidence Interval = certainty of what the mean value is
Confidence interval narrows when more data is added

Distinct from

- Variance measures the width of a distribution
- Standard deviation is $\text{sqrt}(\text{variance})$ and has same units as variable

Variance and Standard Deviation do not change as the population size increases



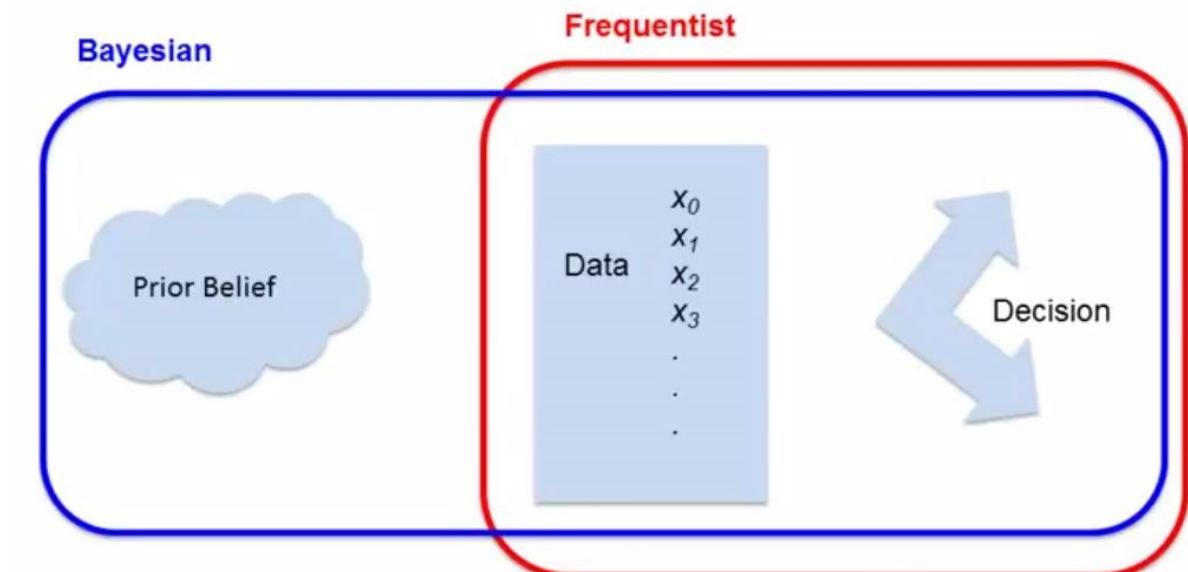
Do we have to rely on experiments?

- What if the system being studied is complex?
- Expensive to replicate?

Bayesian versus Frequentist inference

- *Frequentist* approach measures repeated events and does not depend on a subjective prior that may vary from one investigator to another.
- *Bayesian inference*: "What is the probability that it recently rained given that it is wet outside?"

Both approaches allow evaluation of evidence about competing hypotheses.

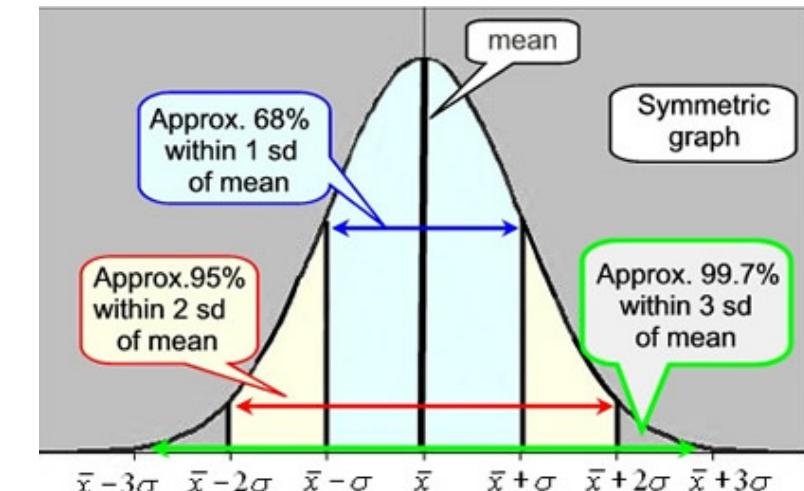


[Jake VanderPlas on difference](#)

- ~~Math and Data Science~~
- ~~Probability~~
- Correlation
- Sets
- Linear Algebra and Numpy
- Calculus and Differential Equations
- Visualization
- Homework

Two events are statistically independent of each other when the probability that one event occurs in no way affects the probability of the other event occurring.

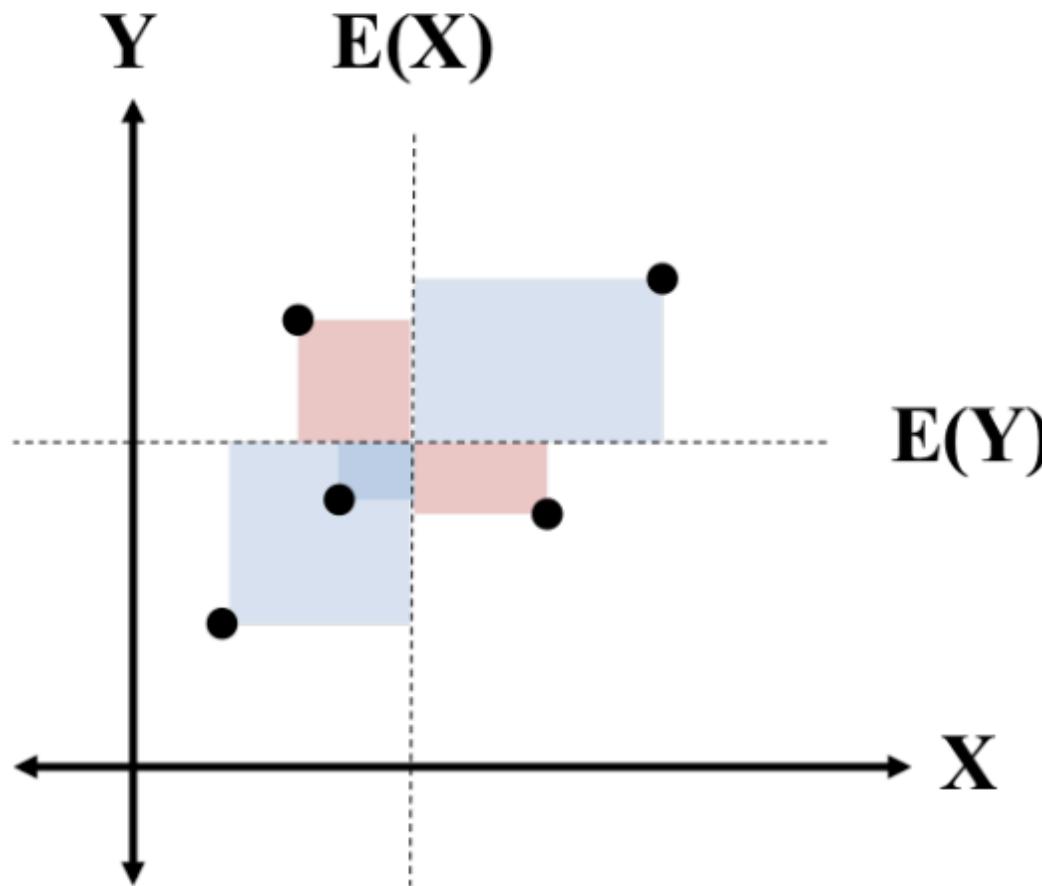
- **Variance** measures width of a distribution
- **Covariance** is the measure of variance for two random variables (joint variability)
- **Correlation** is the normalized covariance, from -1 to 1



More explanation [here](#) and [visualizations of covariance](#) are [helpful](#)

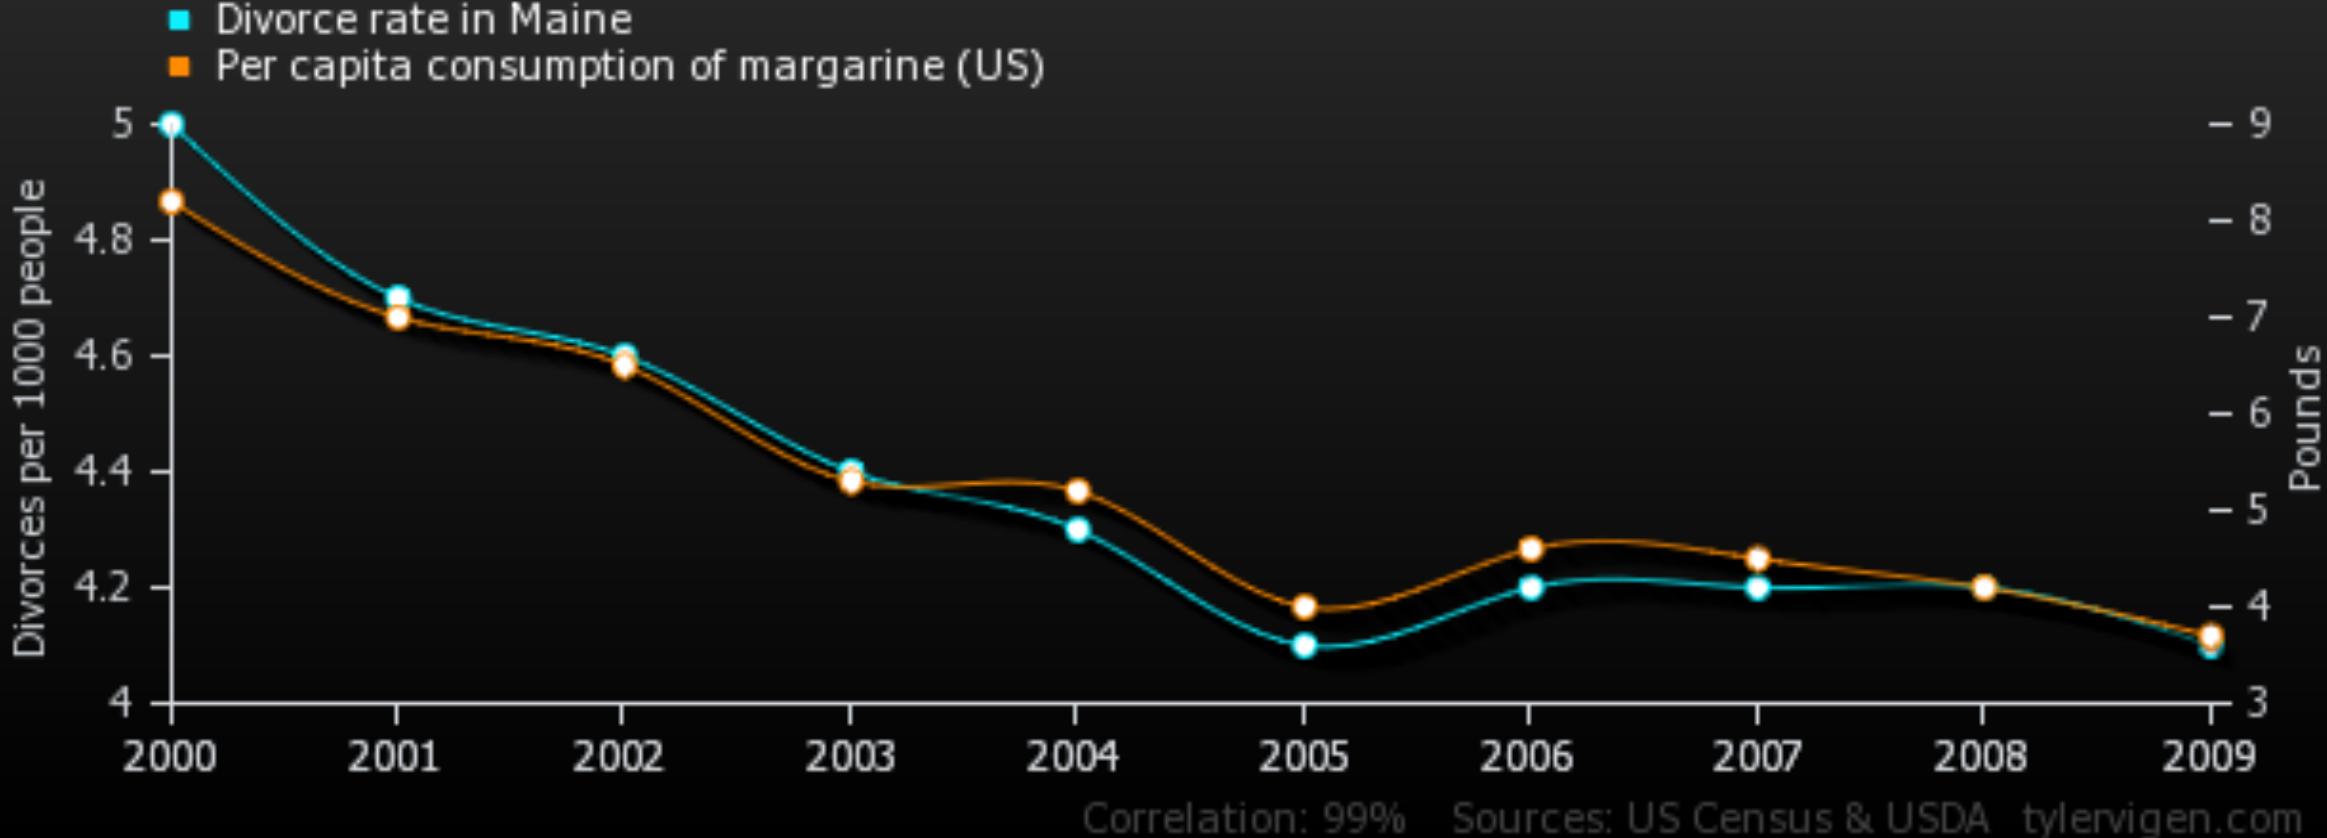
correlation.ipynb

Visualize covariance



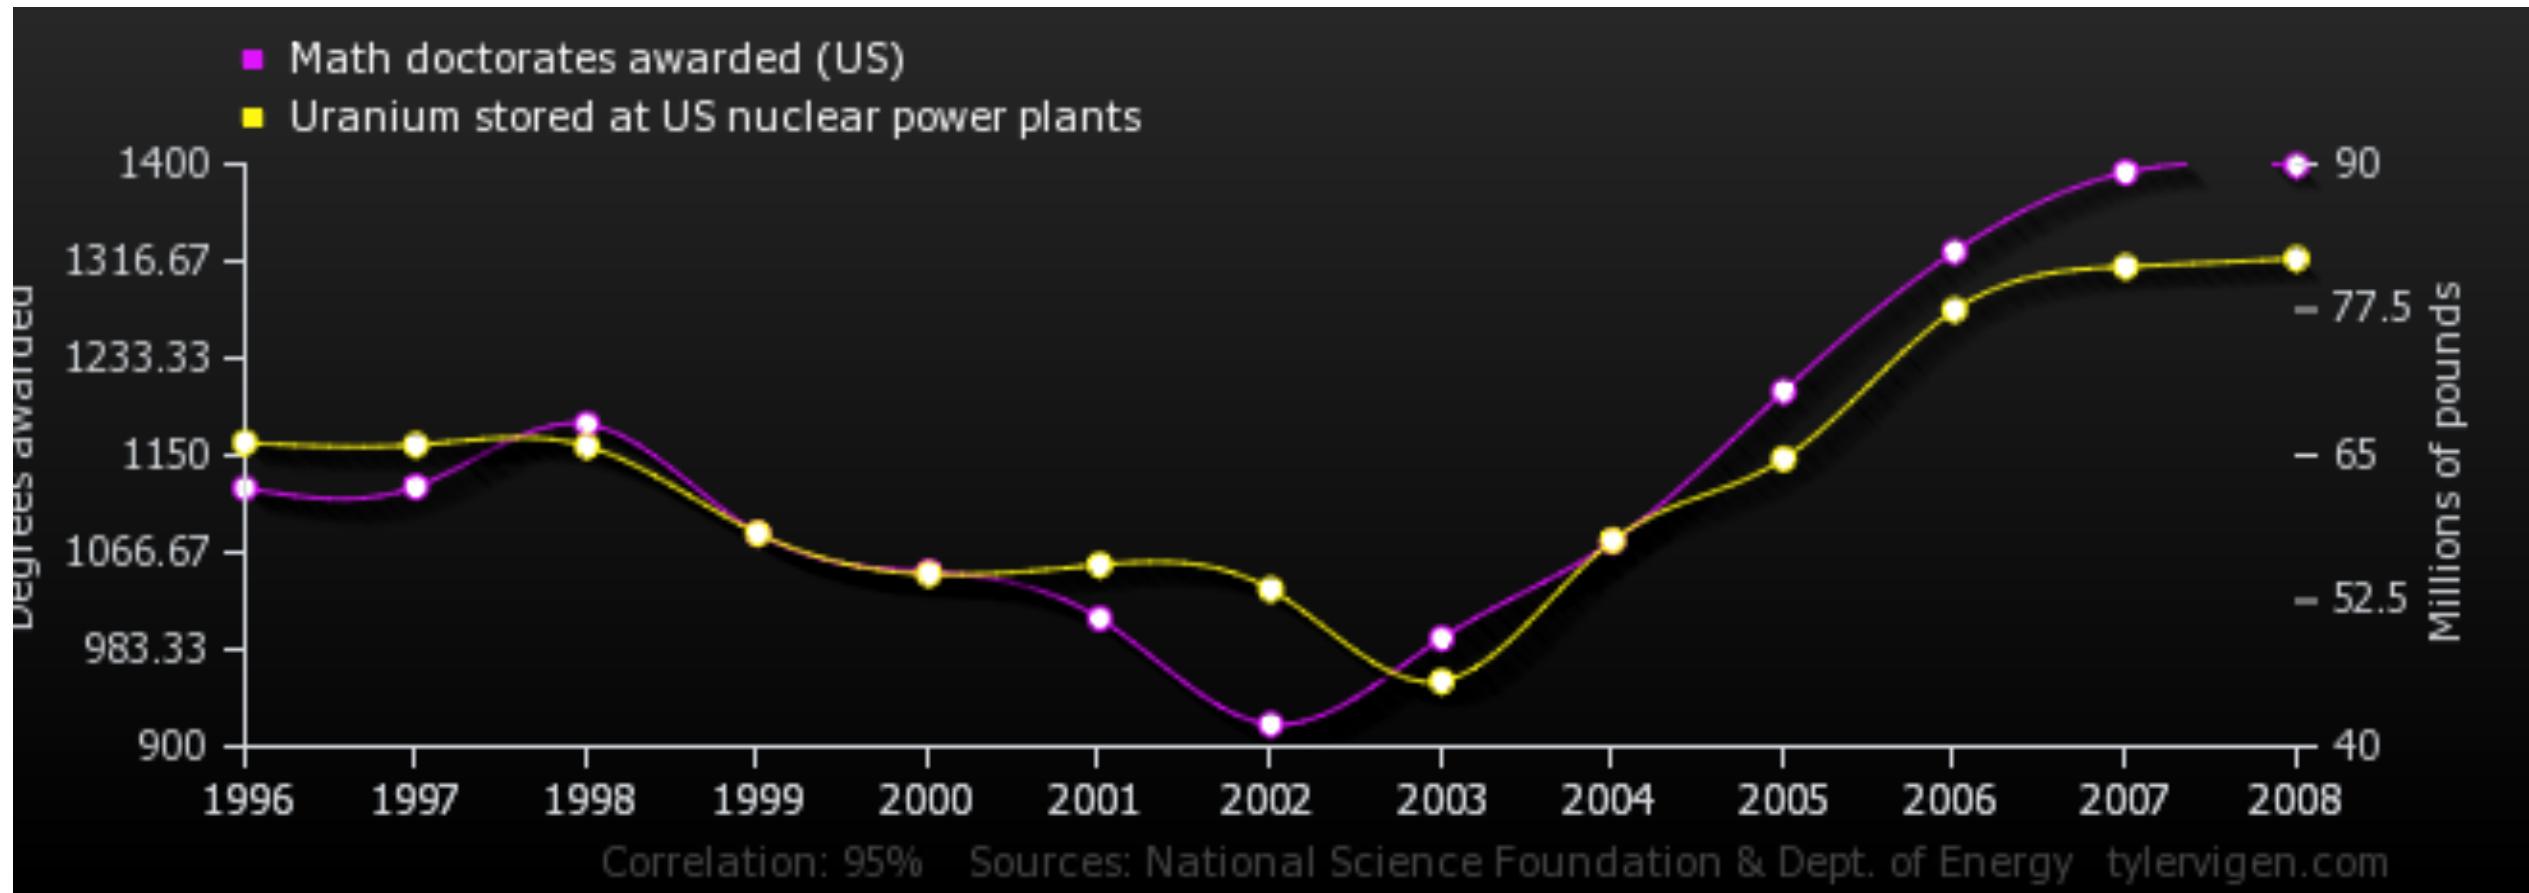
Pink areas are less than the average, so the area has a negative value when summing areas

Correlation = 0.992558 for Divorce and Margarine



- <http://www.tylervigen.com/spurious-correlations>

Correlation = 0.952257 for PhDs and Uranium

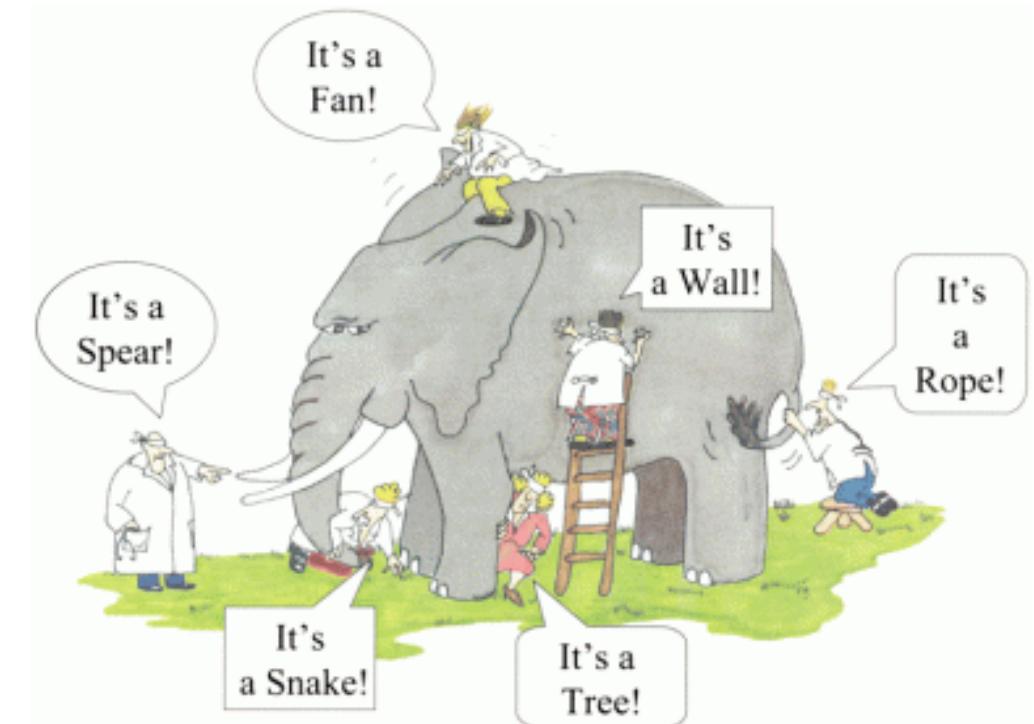


https://tylervigen.com/view_correlation?id=1100

How to (un)intentionally mislead

- Counting and math are objective
- Collection, analysis, and interpretation of data is implemented by humans

Consequence: data you work with may have issues you need to account for



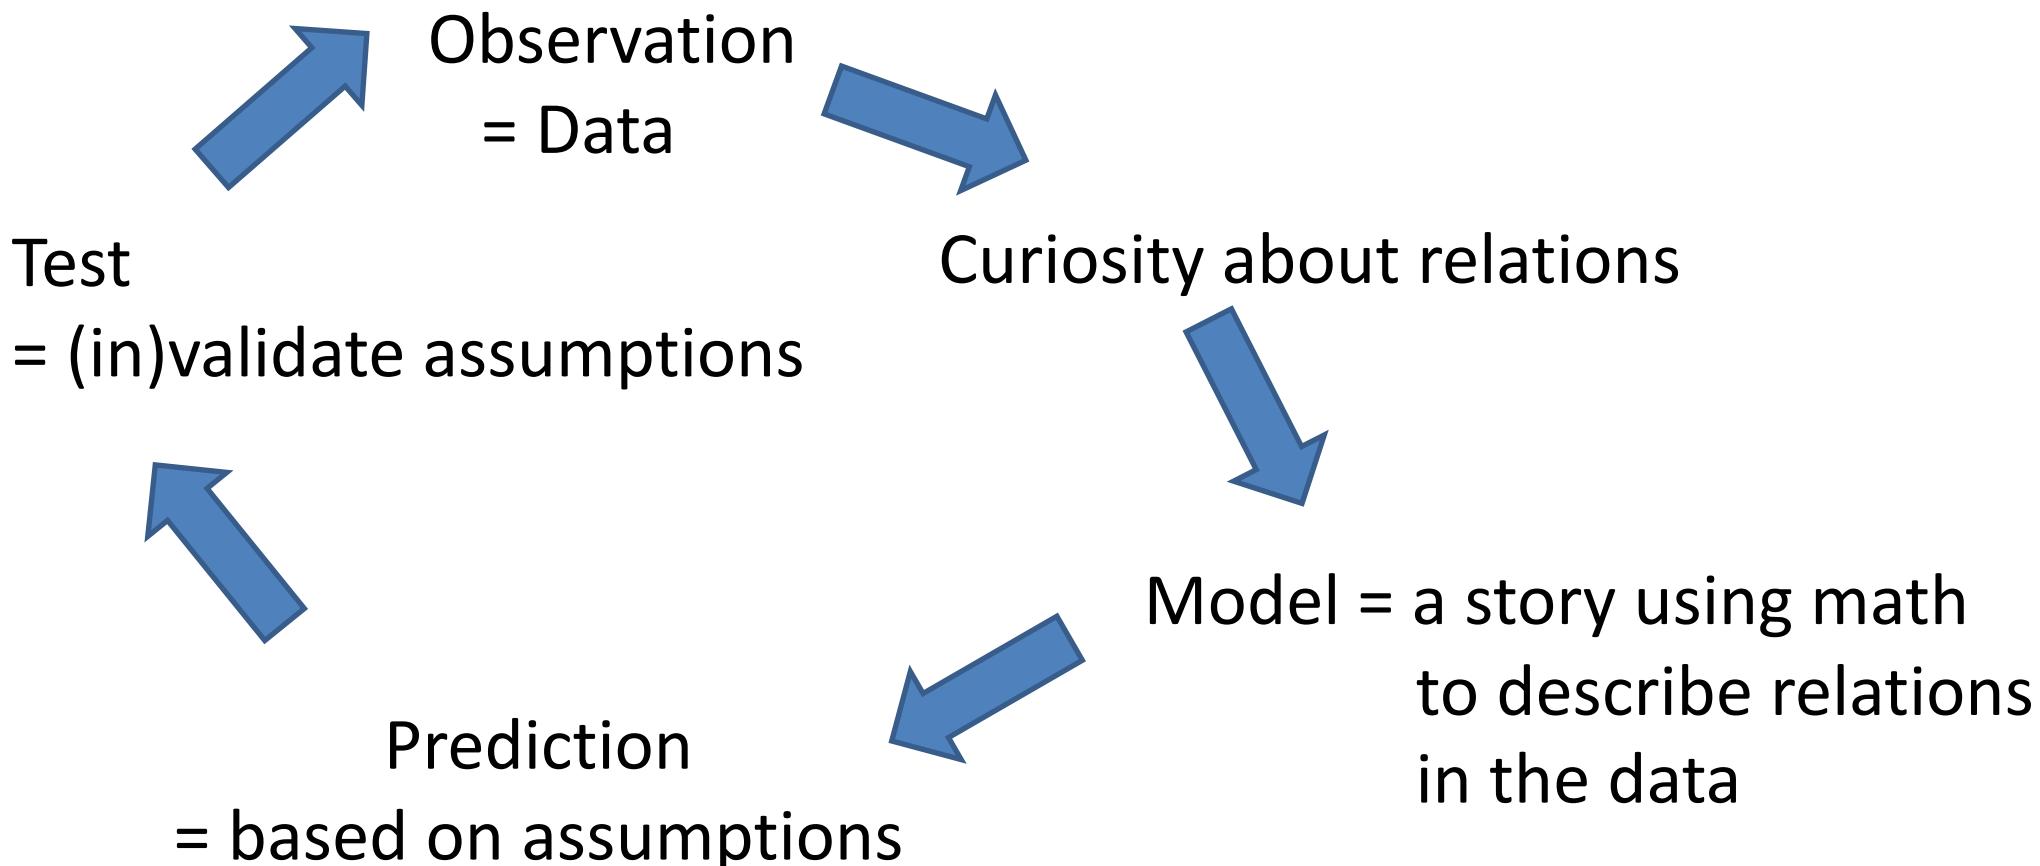
Data collection: exhaust or sample?

Sampling a population can introduce bias

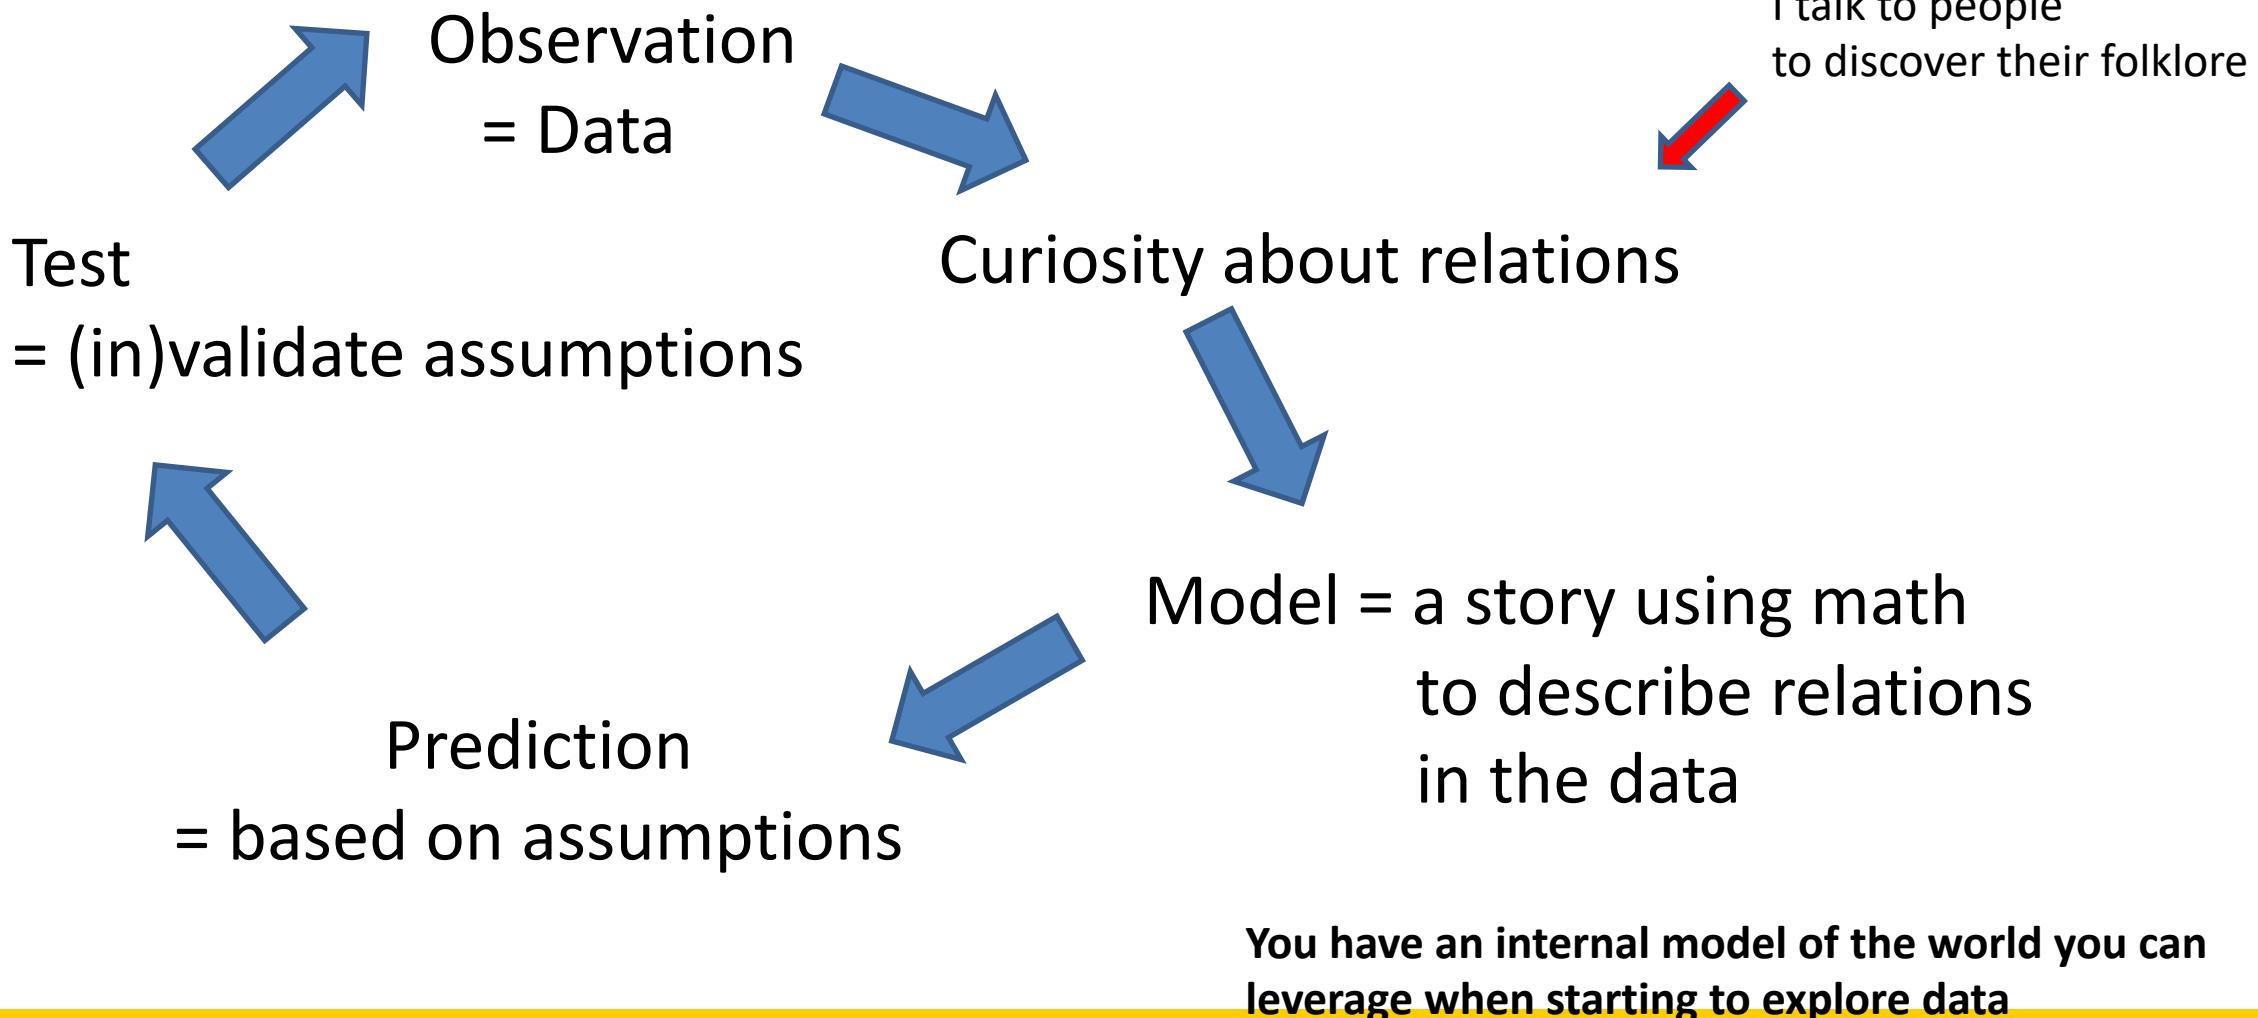
- Area Bias – geographic area of sample needs to be representative of study population
- Self-Selection Bias - decision to participate may correlate with traits that affect the study
- Leading Question Bias - tone of the question suggests the answer
- Social Desirability Bias - reluctance to admit to doing something that is considered socially undesirable

- ~~Math and Data Science~~
- ~~Probability~~
- ~~Correlation~~
- Time series
- Sets
- Linear Algebra and Numpy
- Calculus and Differential Equations
- Visualization
- Homework

Reminder of the Scientific Process



Relevance of folklore in Data Science



Activity: sketch your expectation on paper

Axes:

- Time (days)
- Power (megawatts)



Time varying data

- Source: <https://www.bmreports.com/bmrs/?q=demand/rollingsystemdemand>

Notebook:

visualizing_time_variation_v5_final_product_looks_easy.ipynb

Time varying data

- Source: <https://www.bmreports.com/bmrs/?q=demand/rollingsystemdemand>

Notebook:

visualizing_time_variation_v5_final_product_looks_easy.ipynb

The point of this notebook is that we can tell stories about data just by starting with a scatter plot

In the future, we will discuss a mathematical approach, rather than just relying on visual analysis: [Fourier transform](#)

Tutorial: Groupby

- <https://www.kaggle.com/crawford/python-groupby-tutorial>
- <http://wesmckinney.com/blog/groupby-fu-improvements-in-grouping-and-aggregating-data-in-pandas/>

python-groupby-tutorial.ipynb

Browser history for two people

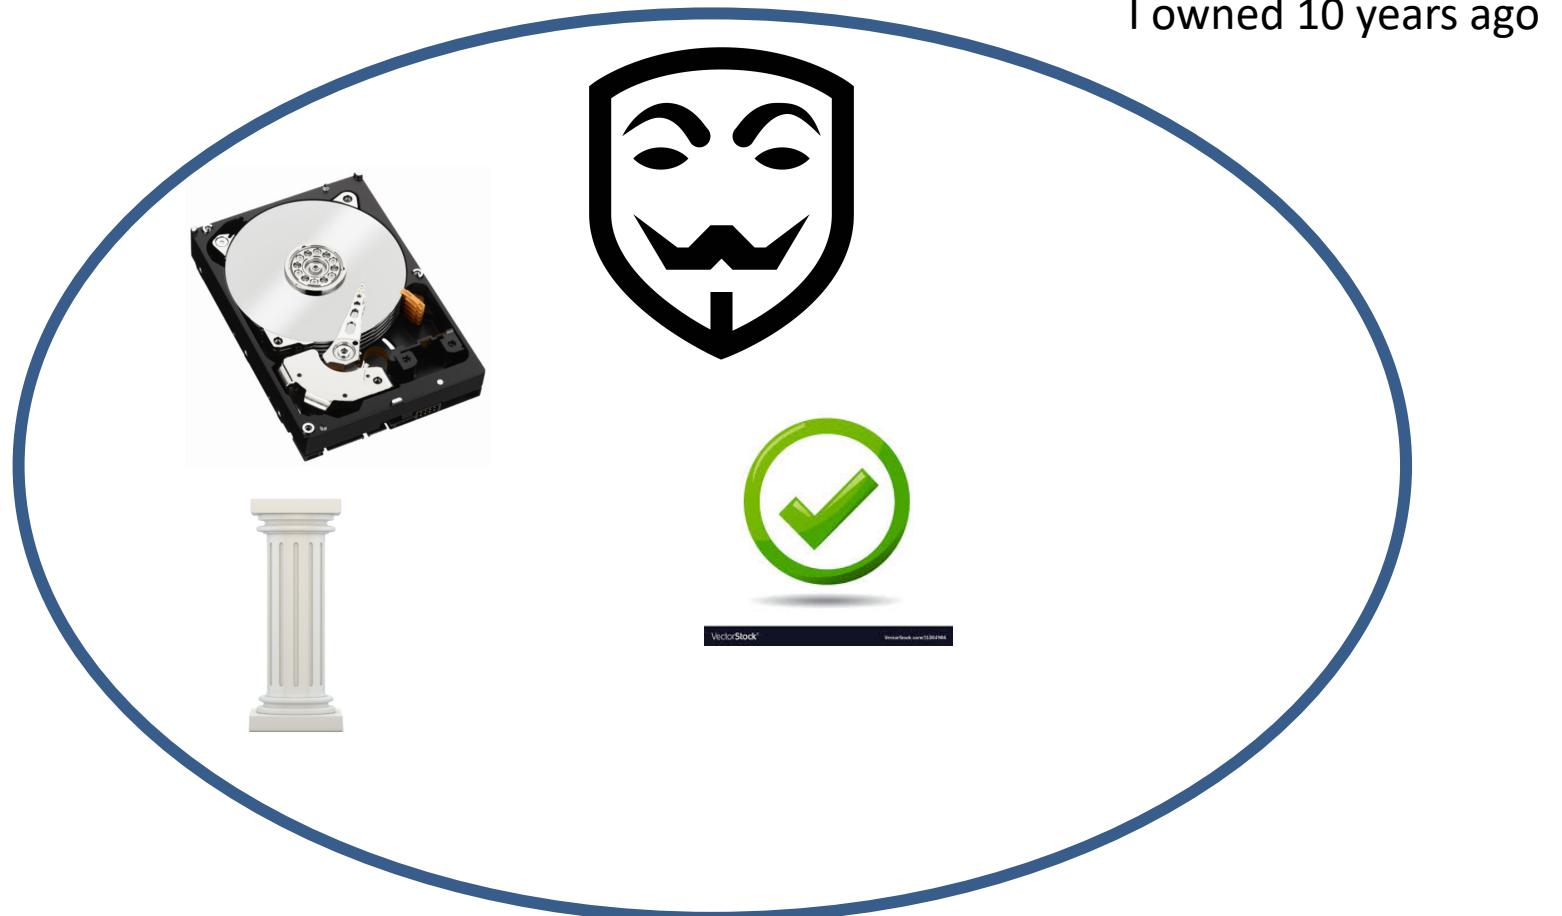
- How would you find which websites both have visited?

Sets are collections of items

Stuff I own now



Sets are collections of items



Sets are collections of items

Stuff I own now



I owned 10 years ago



VectorStock®

Set intersection (and); Set union (or)

- Intersection = the overlap of one or more sets
- Union = combination of one or more sets

sets.ipynb

Linear Algebra and Numpy

- ~~Math and Data Science~~
- ~~Probability~~
- ~~Correlation~~
- ~~Sets~~
- Linear Algebra and Numpy
- Calculus and Differential Equations
- Visualization
- Homework

Relevance of Linear Algebra to Data Science

Central to machine learning

- Images are arrays of numbers
- Text processing (ie [word2vec](#))
- Neural network weights are stored as array
- [Principal Component Analysis](#) (PCA) - lossy dimensionality reduction

<https://www.khanacademy.org/math/linear-algebra>

Numpy and dimensions of data

- Scalar values, ie 5
- 1 dimension
 - Numpy: array – single type of entries
 - Python: list, set (unique elements), tuple (immutable)
 - Linear algebra: vector
- 2 dimensions:
 - CSV: Table; Excel: worksheet
 - Numpy: array of arrays
 - Linear algebra: matrix

NumPy is short for Numerical Python

- ~~Math and Data Science~~
- ~~Probability~~
- ~~Correlation~~
- ~~Sets~~
- ~~Linear Algebra and Numpy~~
- **Calculus and Differential Equations**
- Visualization
- Homework

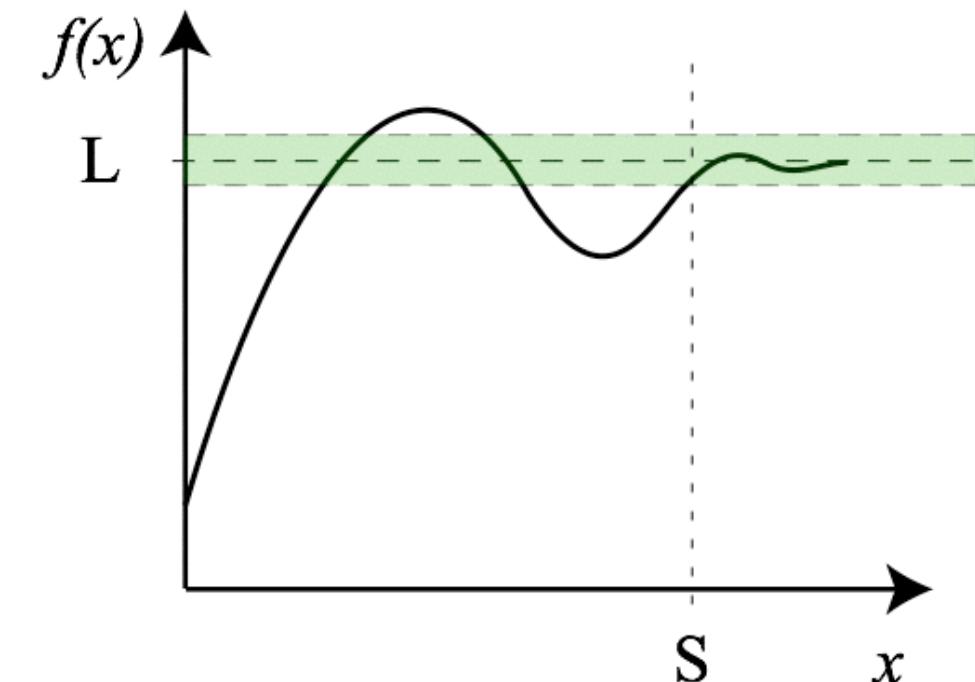
Calculus as gateway to Differential Equations

- Calculus: approximation that summing a large number of small things yields a finite value.

Curious about what calculus covers?

See this [review of concepts](#) and [essential concepts](#).

<https://www.khanacademy.org/math/multivariable-calculus>



Relevance of Differential Equations in Data Science:

Differential equations enable determination of rate of change of quantities



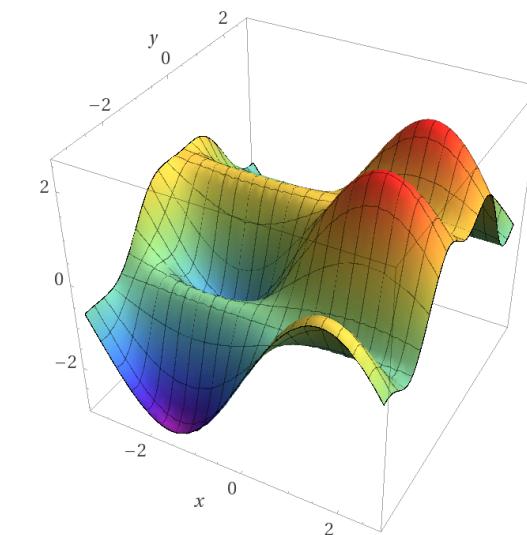
I HAS BOREDOM.

Boredom. I has it.

(When the rate of change is zero,
you are bored)

Relevance of Differential Equations in Data Science

- Differential equations enable determination of rate of change of quantities
- Optimization, which is central to machine learning (ie Gradient descent)
- As with most topics in this lecture, there are entire courses and textbooks dedicated to optimization.



Computed by WolframAlpha

There are other applications

What does a data scientist do all day?

- I have autonomy because my employer expects I know what I'm doing
- Work from 9 to 5 M-F, or 10 to 4, or whatever suits my needs
- Read and write emails.
- Go to meetings; run meetings
- Have impromptu discussions with stakeholders
- Write documentation, read papers
- Write code
- Manage data

No day repeats

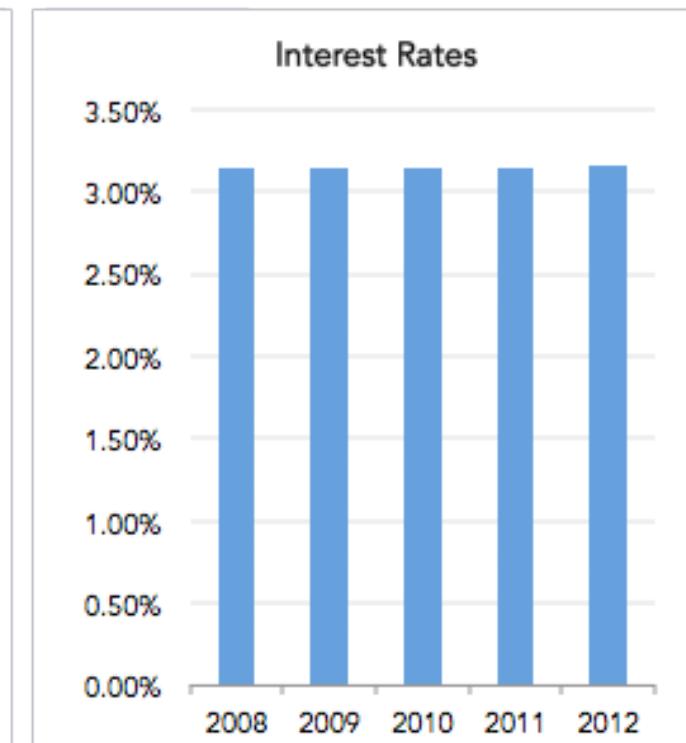
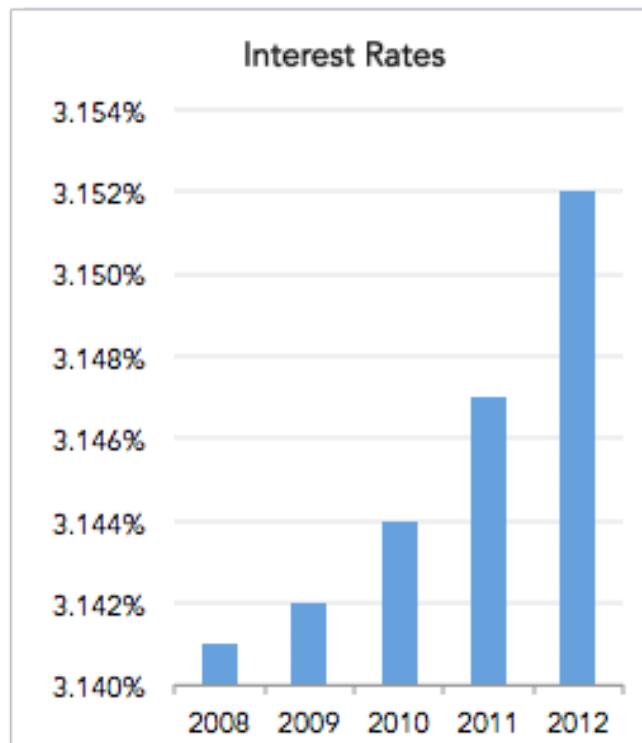
No day is predictable

<https://www.kdnuggets.com/2017/11/day-life-data-scientist.html>

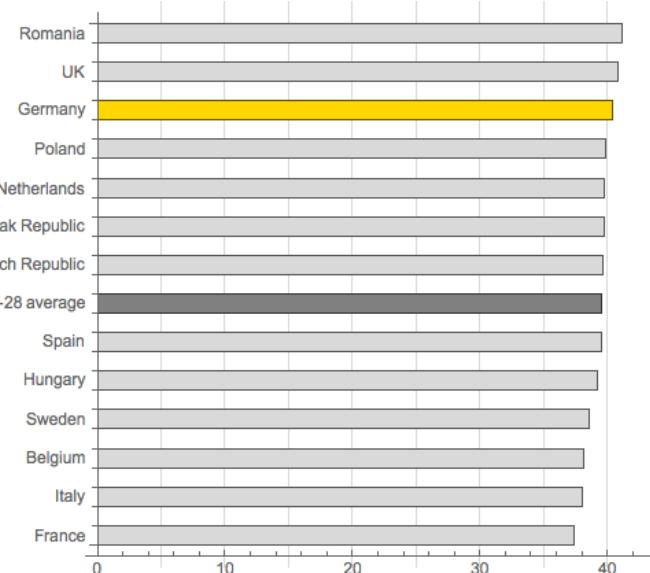
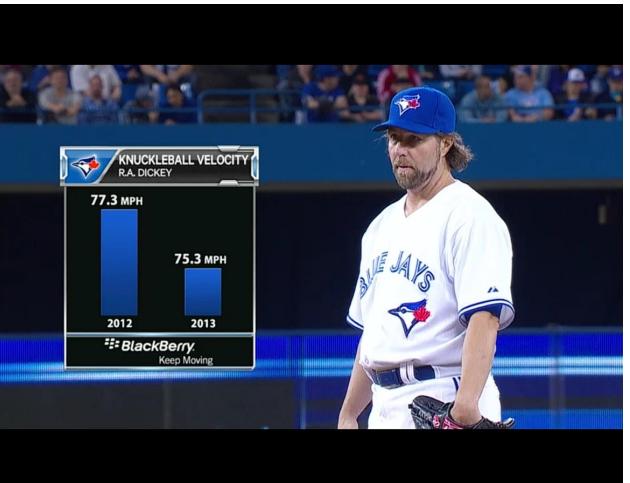
- ~~Math and Data Science~~
- ~~Probability~~
- ~~Correlation~~
- ~~Sets~~
- ~~Linear Algebra and Numpy~~
- ~~Calculus and Differential Equations~~
- **Visualization**
- Homework

Knowing Math insufficient in story telling: Misleading Visualizations

Same Data, Different Y-Axis

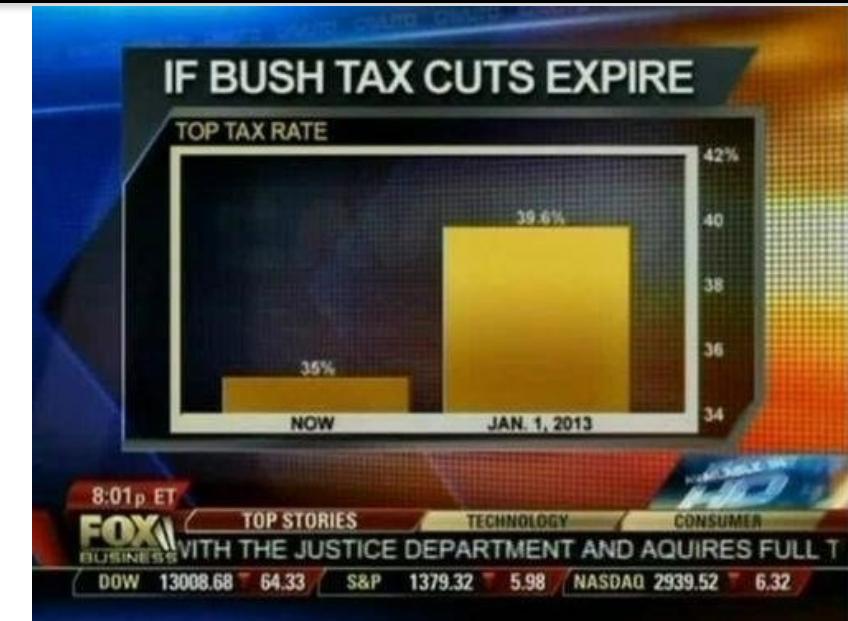
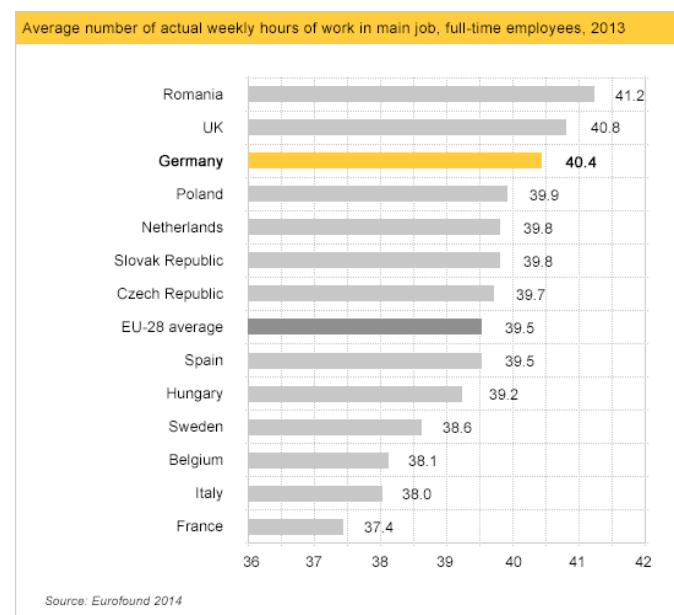


No one would actually do that, right?



Oh yes

Complete data:



(There [are conditions](#) under which not including zero is justified.)

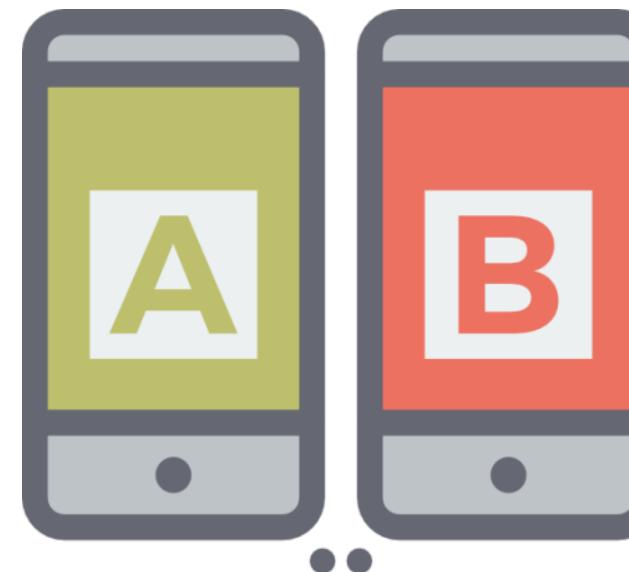
Mistakes when Characterizing data

Even "find the average" isn't as trivial as you may think

- Arithmetic mean : what you're used to
- Harmonic mean : combining multiple parameters which have different ranges so that a given percentage change in any of the properties has the same effect
- Geometric mean : appropriate when the average of rates is desired

There's more than Exploratory Data Analysis

- Comparing models using A/B testing



Where does A/B testing get used?

- Website design
- Email marketing

Useful when the audience being evaluated is sufficiently large to be statistically meaningful

Typically evaluating response with respect to an outcome

Homework

- *Assignment 1:* Acquire power data ([source](#)) for at least 10 days and not more than 40 days. Load the data into a Jupyter Notebook. Create two bar graphs of the power consumption per hour. One bar graph has 24 bars; one bar graph has 24 times number of days bars. Submit the .ipynb file containing the analysis and the generated graphs.
- *Assignment 2:* Simulate a fair die and a biased 6-sided die. The biased die has probabilities $\{0.15, 0.15, 0.15, 0.15, 0.15, 0.25\}$. Create a visualization that compares outcomes of multiple rolls of a fair die and this biased die. You can use a single visualization or multiple visualizations to demonstrate the difference in outcomes for the dice. The user of your notebook should be able to alter the number of simulations as an argument to a function.
[Help on visualizations](#)

